

Machine Learning for Speech Processing

Dr. Mathew Magimai Doss

October 28, 2022

Outline

Introduction

Static classification

Sequence classification

Detection

Outline

Introduction

Static classification

Sequence classification

Detection

Machine learning

- Learning is an essential part of living
- Learning typically means changing/adapting to be better, as per a given criterion, when a similar situation arrives
- Challenge lies in generalizing to new or unobserved situations

Challenge

- Learning needs training data. We have access to only a finite amount of training data.
- Variabilities in the data
 - Language level: isolated words, sentence, spoken language, read speech, spontaneous speech, dialect ...
 - Speaker level: gender, adult versus child, dialect, age, accent, impaired versus unimpaired (pathological speech), emotion, mood, stress
 - Noise
 - Convulsive: recording/transmission condition, reverberation
 - Additive: recording environment, transmission
 - Lombard effect: speaker level variability in noisy environment
- Depending upon the task, a few variabilities are desirable or of interest while others are undesirable or of no interest.

Types of learning

1. Supervised learning

- Training data is labelled. For example, for a frame of feature vector or a sequence feature vectors we have a "class" label associated to it.

2. Reinforcement learning

- Training data has partial labels/targets. For example, did a robot carry out the desired action or not?

3. Unsupervised learning

- Training data does not contain class labels or targets. But, often there is a hidden goal associated with the task. For example, data clustering tasks have a hidden goal such as, minimization of a distance function or maximization of likelihood.

Statistical Pattern Recognition

- Classification
 - Static classification
 - Sequence classification
- Detection: can be regarded as a two class classification problem
- Regression: relation between two variables, namely, measured variable and explanatory variable

Three Key Statistical Rules

1. Bayes's rule:

$$P(A, B) = P(A|B)P(B) = P(B|A)P(A)$$


2. If B_k ($k = 1, \dots, K$) are mutually exclusive and collectively exhaustive ($\sum_{k=1}^K P(B_k) = 1$)

$$P(A) = \sum_{k=1}^K P(A, B_k)$$

3. Gibbs sampler:

$$P(B_1, \dots, B_k, \dots, B_K) = \prod_{k=1}^K P(B_k | B_{k-1}, \dots, B_1)$$


Static Classification



Sequence classification



Detection



Outline


Introduction

Static classification

Sequence classification

Detection

Static Classification



Theoretical formulation

$$P(C_k|x_m, \Theta) = \frac{p(x_m|C_k, \Theta) \cdot P(C_k|\Theta)}{p(x_m|\Theta)} \quad \forall k \in \{1, \dots, K\} \quad (1)$$

- $P(C_k|x_m, \Theta)$: Posterior probability of class C_k
- $p(x_m|C_k, \Theta)$: Likelihood of class C_k
- $P(C_k|\Theta)$: Prior probability of class C_k
- $p(x_m|\Theta) = \sum_{j=1}^K p(x_m|C_j, \Theta) \cdot P(C_j|\Theta)$: Observation likelihood
- Θ : parameters of the statistical model
- $0 \leq P(C_k|x_m, \Theta) \leq 1$ and $\sum_{k=1}^K P(C_k|x_m, \Theta) = 1$
- $0 \leq P(C_k|\Theta) \leq 1$ and $\sum_{k=1}^K P(C_k|\Theta) = 1$

Generative approach

- Estimate or model $p(x_m | C_k)$ by a probability density function
 - Gaussian or Normal distribution

$$\begin{aligned}
 p(x_m | C_k, \Theta_k) &= N(x_m, \mu_k, \Sigma_k) \\
 &= \frac{1}{(2\pi)^D |\Sigma_k|^{1/2}} \exp\left(-\frac{(x_m - \mu_k)^t \Sigma_k^{-1} (x_m - \mu_k)}{2}\right) \\
 &\approx \prod_{d=1}^D \frac{1}{\sqrt{2\pi} \sigma_k^d} \exp\left(-\frac{1}{2} \left(\frac{x_m^d - \mu_k^d}{\sigma_k^d}\right)^2\right)
 \end{aligned}$$

- Gaussian mixture models (GMM)

$$p(x_m | C_k, \Theta_k) = \sum_{j=1}^J c_k^j \cdot N(x_m, \mu_k^j, \Sigma_k^j)$$

- Estimate prior probability $P(C_k)$ (typically done through counting)
- Apply Eqn. (1)
- Θ : means, variance and Gaussian weights (in the case of GMMs)

EM algorithm for GMMs (1)

$$p(x_m) = \sum_{j=1}^J c_j p(x_m | \mu_j, \Sigma_j)$$

with $c_j = P(G_j)$ (Weight for Gaussian j).

Estimation step:

$$P(G_j | x_m) = \frac{c_j p(x_m | \mu_j^{(t)}, \Sigma_j^{(t)})}{\sum_j c_j p(x_m | \mu_j^{(t)}, \Sigma_j^{(t)})}$$


Maximization step:

$$\mu_j^{(t+1)} = \frac{\sum_{m=1}^{M} x_m P(G_j | x_m)}{\sum_{m=1}^{M} P(G_j | x_m)}$$

$$\Sigma_j^{(t+1)} = \frac{\sum_{m=1}^{M} P(G_j | x_m) (x_m - \mu_j^{(t+1)}) (x_m - \mu_j^{(t+1)})^T}{\sum_{m=1}^{M} P(G_j | x_m)}$$


$$c_j^{(t+1)} = \frac{1}{M} \sum_{m=1}^{M} P(G_j | x_m)$$

EM algorithm for GMMs (2)




Discriminative approach

- Artificial neural network trained with one hot encoding of target and cross entropy error function (or mean square error function) can directly estimate $P(C_k|x_m, \Theta)$



- Support vector machines (estimation of posterior probability of class not trivial, see **Platt's method**)
- Θ : parameters of the artificial neural networks (weights and biases) or support vector machines

ANN: Multilayer perceptrons



Typical multilayer perceptron (MLP) architecture, each unit approximating a perceptron.

ANN Training

Supervised training: Input vector sequence:

$$X = \{x_1, \dots, x_n, \dots, x_M\}$$

Desired output sequence associated with X :

$$D = \{d(x_1), \dots, d(x_m), \dots, d(x_M)\}$$

and $d(x_m) = (d_1(x_m), \dots, d_k(x_m), \dots, d_K(x_m))^T$

In classification mode:

$$d_k(x_m) = \delta_{k\ell} \text{ if } x_m \in C_\ell$$

Parameters Θ : weights and biases

Training Criteria

- Mean Square Error:

$$\operatorname{argmin}_{\{\Theta\}} E = \frac{1}{2} \sum_{m=1}^M \sum_{k=1}^K [g_k(x_m, \Theta) - d_k(x_m)]^2$$

- Entropy or relative entropy:

$$\begin{aligned} \operatorname{argmin}_{\{\Theta\}} E_e &= \sum_{m=1}^M \sum_{k=1}^M \left[d_k(x_m) \log \frac{d_k(x_m)}{g_k(x_m, \Theta)} \right. \\ &\quad \left. + (1 - d_k(x_m)) \log \frac{1 - d_k(x_m)}{1 - g_k(x_m, \Theta)} \right] \end{aligned}$$

$g_k(x_m, \Theta)$ denotes the output of the neural network.

Error back propagation training

Minimization of E (or E_e) in the parameter space Θ (weights + biases)

$$\frac{\nabla E}{\nabla \Theta} = 0$$

Done via a gradient procedure:

$$\Delta w_{ij} = -\alpha \frac{\partial E}{\partial w_{ij}}$$

α denotes learning rate

Adjust w_{ij} based on Δw_{ij} .

Offline Error Back-Propagation

Initialize network at random; choose “large” learning rate

Until convergence = true

For m=1 to M

 Forward computation of $g_k(x_m, \Theta)$

 Error calculation and global error update

 Error backward propagation

 and compute local $\delta\Theta(x_m)$

Θ update = $\sum_{m=1}^M \delta\Theta(x_m)$

If error (on cross-validation set) decreases

 save new parameters

Otherwise, don't save new parameters

 and decrease learning rate

If learning rate < threshold then convergence = true;

Online Error Back-Propagation

Initialize network at random

Choose “large” learning rate; convergence = false

Until convergence = true

 For m=1 to M (or something else)

 Pick x_m at random

 Forward computation of $g_k(x_m, \Theta)$

 Error calculation

 Error backward propagation and Θ update

 If error (on cross-validation set) decreases

 save new parameters

 Otherwise, don't save new parameters

 and decrease learning rate

If learning rate < threshold: convergence = true;

In practice: Mini-batch training, a combination of offline and online error back propagation.

Cross-validation training (1)


After each MLP training epoch:

1. check recognition performance on independent data set
2. stop training if rec performance starts to decrease and learning rate below a given threshold

Remarks: there are other solutions like


- Forcing small weights
- “Optimal Brain Damage”
- Regularized training (Bayesian approach)

Cross-validation training (2)



Example of crossvalidation training. x-val represents the crossvalidation data on which classification performance is regularly checked. Training is stopped when performance on x-val data reaches the maximum.

Interpretation of ANN output $g_k(x_m, \Theta)$



ANN output $g_k(x_m, \Theta)$ is an estimate of posterior probability of class C_k , i.e. $P(C_k|x_m)$.

Decision making

- Maximum likelihood

$$C_{mle}^* = \arg \max_k p(X|C_k, \Theta)$$

$$p(X|C_k, \Theta) = \prod_{m=1}^M P(x_m|C_k, \Theta) \text{(Assuming i.i.d)}$$

- Maximum a posteriori probability

$$C_{map}^* = \arg \max_k p(C_k|X, \Theta)$$

$$P(C_k|X, \Theta) = \frac{1}{Z} \cdot \prod_{m=1}^M P(C_k|x_m, \Theta) \text{(Assuming i.i.d)}$$

or

$$P(C_k|X, \Theta) = \frac{1}{Z} \cdot \sum_{m=1}^M P(C_k|x_m, \Theta)$$

Z is a normalization factor.

Better to perform computation using logarithm to avoid underflow issues.

Outline

Introduction

Static classification

Sequence classification

Detection

Sequence classification



Statistical ASR

Given feature observation $X = \{x_1, \dots, x_m, \dots, x_M\}$ predict the most probable word sequence $W^* = \{w_1^* \dots w_n^* \dots w_N^*\}$

$$\begin{aligned} W^* &= \arg \max_{W_i \in \mathcal{W}} P(W_i | X, \Theta) \\ &= \arg \max_{W_i \in \mathcal{W}} \frac{p(X | W_i, \Theta_a) \cdot P(W_i | \Theta_l)}{p(X | \Theta)} \\ &= \arg \max_{W_i \in \mathcal{W}} p(X | W_i, \Theta_a) \cdot P(W_i | \Theta_l), \end{aligned}$$

where W_i denotes a word hypothesis, \mathcal{W} denotes a set of word hypotheses and $\Theta = \{\Theta_a, \Theta_l\}$

- Acoustic modeling: estimation of $p(X | W_i, \Theta_a)$ using hidden Markov models (HMMs)
- Language modeling: estimation of $P(W_i | \Theta_l)$ using discrete Markov models (DMMs)

Discrete Markov Model (DMM)


- Stochastic finite state automaton
- MM built up from states q_ℓ from a set of classes (states)
 $\Omega = \{\omega_1, \dots, \omega_k, \dots, \omega_K\}$
- q^n particular state of M visited at time n ,
- $q_\ell^n \equiv \{q^n = q_\ell\}, q_\ell \in \Omega$
- MM is defined by topology, i.e., how the states are connected
See HMM lab exercise 1
- Parametrized by:

$$\begin{aligned} P(q_\ell^n | q_k^{n-1}, q_j^{n-2}, \dots) &\simeq P(q_\ell^n | q_k^{n-1}) \quad (1st \text{ order Markov}) \\ &\simeq P(q_\ell | q_k) = P_{k\ell} \quad (\text{time independent}) \end{aligned}$$

Note: $q_\ell, q_k \in \Omega$

Transition probability matrix: $A = \{P_{k,\ell}\}$.

DMM (2)



Example of fully connected discrete Markov model with $\Omega = \{a, b, c\}$. For example, in case of weather model: “a” = “cloudy”, “b”= “rainy” and “c” = “sunny”

Typical Problems (1)

■ Probability of a particular path

$$\begin{aligned} P(Q|MM) &= P(q^1|q_i^0)P(q^2|q^1)\dots P(q^n|q^{n-1})\dots P(q^N|q^{N-1}) \\ &= \prod_{n=1}^N P(q^n|q^{n-1}) \end{aligned}$$

One way to estimate $P(W_i|\Theta_I)$, e.g.

$$P(W^*|\Theta_I) = \prod_{n=1}^N P(w_n^*|w_{n-1}^*)$$

■ State duration distribution

Probability to stay in state q_i for *exactly* d time steps?

$$Q = \{q_i^0, q_i^1, q_i^2, \dots, q_i^d, q_j^{d+1}\}, \text{ with } j \neq i$$

and:

$$P(Q|MM) = (P_{ii})^{d-1}(1 - P_{ii})$$

See HMM lab exercise 1

Typical Problems (2)

- Probability to go from state q_i to q_j in N steps

$$P(q_j^N | q_i^0) = \sum_{n=0}^N \sum_{\ell=1}^L P(q_j^N, q_\ell^n | q_i^0)$$

Defining

$$\alpha(\ell, n) = P(q_\ell^n | q_i^0, N)$$

We have:

$$\alpha(\ell, n+1) = \sum_k \alpha(k, n) P_{k\ell}$$

$$P(q_j^N | q_i^0) = \alpha(j, N)$$

Typical Problems (3)

- **Probability of best path of length N between q_i and q_j**
If $\bar{P}(k, n)$ is probability of best path to go from q_i to q_k in n steps:

$$\bar{P}(\ell, n+1) = \max_k \bar{P}(k, n) P_{k\ell}$$

and

$$\bar{P}(q_j^N | q_i^0) = P(j, N)$$

Generalization:

$$A^n(i, j) = P(q_j^n | q_i^0)$$


DMM training

- Training data
 - Seq 1: $R, R, R, S, S, C, R, C, \dots$
 - Seq 2: S, C, S, C, C, R, \dots
 - .
 - .
 - Seq L: $S, R, C, R, C, R, C, S, S, \dots$
- $P(S/I) = \text{Count(seq starting with } S\text{)} / L$, $P(R/I) = \text{Count(seq starting with } R\text{)} / L$ (L denotes number of sequences in the training data)
- $P(S/R) = \text{Count}(R \rightarrow S) / \text{Count}(R)$, $P(C/S) = \text{Count}(S \rightarrow C) / \text{Count}(S)$, $P(R/R) = \text{Count}(R \rightarrow R) / \text{Count}(R)$

Another approach: train a neural network, e.g., recurrent neural network


From DMM to HMM (1)

- Coin tossing and only the results of each coin toss i.e. Heads (H) or Tails (T) is revealed
 $H\ H\ T\ T\ T\ H\ H\ H\ H\ T\ T\dots$
- 1-coin model: DMM with two states $\Omega = \{H, T\}$ and all transition probabilities are equal to 0.5.



From DMM to HMM (2)

- 2-coin model: HMM with two states $\Omega = \{C_1, C_2\}$; each state associated with $P(H|C_i)$ and $P(T|C_i)$ $i \in \{1, 2\}$; and the transition probability reflects the probability of choosing one of the (possibly biased) coins



Estimation of $P(X|W_i, \Theta_a)$ using HMMs

$$\begin{aligned} P(X|W_i, \Theta_a) &= \sum_{Q \in W_i} P(X, Q|W_i, \Theta_a) \\ &= \sum_{Q \in W_i} P(X|Q, W_i)P(Q|W_i), \end{aligned}$$

where $Q = \{q_1, \dots, q_m, \dots, q_M\}$ denotes sequence of HMM states.

After i.i.d and first order Markov assumption

$$\begin{aligned} P(X|W_i, \Theta_a) &= \sum_{Q \in W_i} \prod_{m=1}^M p(x_m|q_m) \prod_{m=1}^M P(q_m|q_{m-1}) \\ &= \sum_{Q \in W_i} \prod_{m=1}^M p(x_m|q_m)P(q_m|q_{m-1}) \text{ Full likelihood} \\ &\approx \max_{Q \in W_i} \prod_{m=1}^M p(x_m|q_m)P(q_m|q_{m-1}) \text{ Viterbi approx.} \end{aligned}$$

For the sake of simplicity, Θ_a and W_i are dropped in the latter equations.

$p(x_m|q_m)$ is referred to as emission likelihood and $P(q_m|q_{m-1})$ state transition probabilities. See HMM lab exercises 3 and 4.

Estimation of emission likelihood

- Gaussians and Gaussian-Mixtures

$$p(x_m | q_m = k, \Theta_a) = N(x_m, \mu_k, \Sigma_k)$$

$$= \frac{1}{(2\pi)^{D/2} |\Sigma_k|^{1/2}} \exp \left(-\frac{(x_m - \mu_k)^t \Sigma_k^{-1} (x_m - \mu_k)}{2} \right)$$

OR :


$$p(x_m | q_m = k, \Theta_a) = \sum_{j=1}^J c_k^j N(x_m, \mu_k^j, \Sigma_k^j)$$

- Artificial neural networks: estimate scaled-likelihood

$$p_{sl}(x_m | q_m = k, \Theta_a)$$

$$p_{sl}(x_m | q_m = k, \Theta_a) = \frac{p(x_m | q_m = k, \Theta_a)}{p(x_m | \Theta_a)} = \frac{P(q_m = k | x_m, \Theta_a)}{P(q_m = k | \Theta_a)}$$




Hidden Markov Model (HMM)



Double, embedded stochastic process:

- choose box using transition probs.
 - choose ball using emission probs.
- Automatic Speech Processing, Masters Cycle

Hidden Markov Model (HMM)



$$\begin{aligned} P(r, b, r | \text{Model}) &= P(r, b, r | B_1, B_1, B_1) P(B_1, B_1, B_1) \\ &+ P(r, b, r | B_1, B_1, B_2) P(B_1, B_1, B_2) \\ &+ P(r, b, r | B_1, B_1, B_3) P(B_1, B_1, B_3) \\ &+ P(r, b, r | B_1, B_2, B_1) P(B_1, B_2, B_1) \\ &+ P(r, b, r | B_1, B_2, B_2) P(B_1, B_2, B_2) \\ &\quad \dots \\ &+ P(r, b, r | B_3, B_3, B_3) P(B_3, B_3, B_3) \end{aligned}$$

with $P(r, b, r | B_i, B_j, B_k) = P(r | B_i) P(b | B_j) P(r | B_k)$

Conditional i.i.d assumption, i.e., given the state at any time instance the observation at that time instance/frame is only dependent on that state and independent of all observations and states at other time instances/frames

$$P(B_i, B_j, B_k) = P(B_i) P(B_j | B_i) P(B_k | B_j)$$

First order Markov assumption

Hidden Markov Model (HMM)

- **Estimation** problem : $P(X|M_w)$?

$O(N\text{states}^{N\text{obs}})$ computations???

Dynamic programming

Baum-Welch


(all paths)

Viterbi

(best path only)


- **Training** problem: best emission and transition probs, given model topology and data
- **Decoding** problem: best sequence of states
→ **Viterbi** again


Training HMMs



- Data: « brbrrbgbrbbgbggbrbgbrggbrb... »
- Emission and transition probabilities?
If states were known: counting
- **EM** (expectation-maximization) **algorithm**
 - Initialize Probs. (first guess if possible): M_w^0
 - Decode the data with $M_w^0 \rightarrow$ states
 - Re-estimate Probs. by counting: $M_w^1 \rightarrow M_w^j \approx M_w^{j+1}$ until



Word sequence modeling (DMM)







Each phone is modeled by an HMM

- Observation= $x_m \Rightarrow p(x_m | \text{phone state}) = \text{continuous}$
Cannot be estimated by counting
Estimated via the p.d.f. of a distribution
 - ex: Gaussian
 - Multi-Gaussian



- In practice, each phoneme is modeled as 3 states (*Bakis* model)





Outline

Introduction

Static classification

Sequence classification

Detection

Detection



$$\frac{p(X|\text{genuine speaker})}{p(X|\text{impostor})} \geq \delta_{\text{asv}}$$

Challenge: modeling or estimation of $p(X|\text{impostor})$



$$\frac{p(X|\text{keyword})}{p(X|\text{not keyword})} \geq \delta_{\text{kws}}$$

Challenge: modeling or estimation of $p(X|\text{not keyword})$

Two types of error: False negative, False positive

Thank you for your attention!

Dr. Mathew Magimai Doss

Idiap Research Institute, Martigny, Switzerland